

## Assignment 1, Cloud Computing

24MD0470 - Akhmetkan Azhar

### Exercise 1: Understanding Cloud Computing Models

- **Objective:** Explore different cloud computing models and understand their key differences.
- **Steps:**
  - Research the three primary cloud service models: Infrastructure as a Service (IaaS), Platform as a Service (PaaS), and Software as a Service (SaaS).
  - Create a table comparing these models in terms of control, flexibility, and use cases.
  - Identify examples of services offered by Google Cloud Platform (GCP) under each model.

- **Questions:**

- What are the main differences between IaaS, PaaS, and SaaS?

**IaaS or Infrastructure as a Service** provides virtualized computing resources over the internet and is particularly useful for custom hardware setups. Meanwhile, **PaaS or Platform as a Service** allows users to manage applications without having to address the underlying infrastructure. PaaS is perfect for quickly building and deploying apps. On the other hand, **SaaS or Software as a Service** requires a subscription to access software features over the internet. SaaS offers the least control, as the provider manages everything, making it ideal for collaboration tools.

- Which GCP services fall under each of these models?

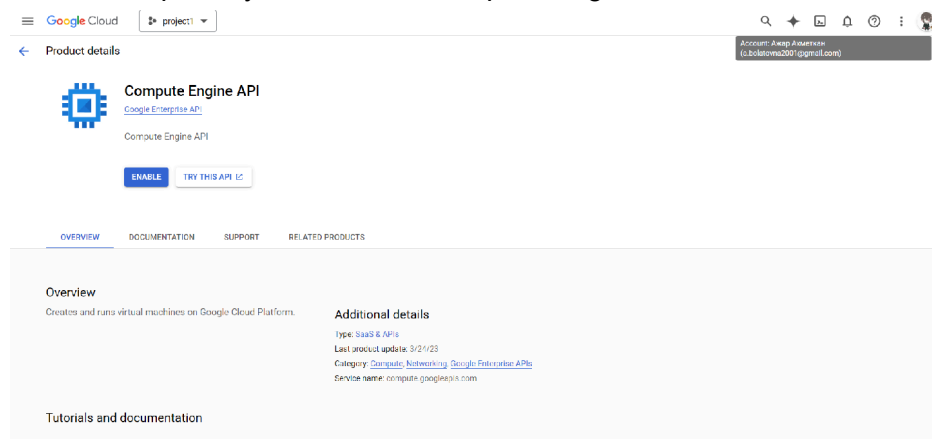
In the context of Google Cloud Platform, these models are well represented. For **IaaS**, services such as Google Compute Engine and Google Cloud Storage are prominent. In the **PaaS** category, Google App Engine and Google Cloud Functions cater to application development needs. In the **SaaS** realm, Google Workspace—including tools like Gmail and Google Docs—provides convenient solutions for everyday tasks.

- Provide a real-world example where each cloud service model might be the most appropriate choice.

Lets illustrate these concepts with real-world examples. If someone starts a new company and needed a custom server for an application then they would likely choose Google Compute Engine. Google App Engine would be the best option for working on a web app who want to focus on coding instead of server management. And for daily tasks like managing emails and collaborating with a team, subscribing to Google Workspace would be a best choice, as it requires no maintenance from the user.

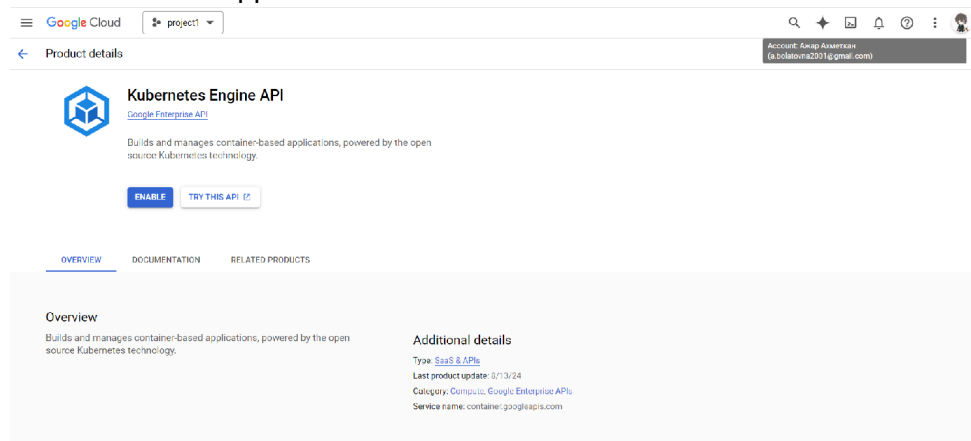
### Exercise 2: Exploring Google Cloud Platform's Core Services

- **Objective:** Get acquainted with the core services provided by Google Cloud Platform.
- **Steps:**
  - Access the Google Cloud Console and navigate to the list of GCP services.
  - Explore and describe the purpose of the following core services:
    - Compute Engine
    - Google Kubernetes Engine (GKE)
    - App Engine
    - Cloud Storage
    - BigQuery
  - For each service, identify a potential use case in a business scenario.
- **Questions:**
  - What is the primary use case of Compute Engine?



the primary use case of Compute Engine is to offer scalable and customizable virtual machines for different tasks. Businesses can use it to run applications, host websites, or handle batch processing. It's especially helpful for companies that want the flexibility to adjust their virtual machines according to specific needs, like CPU, memory, and storage options.

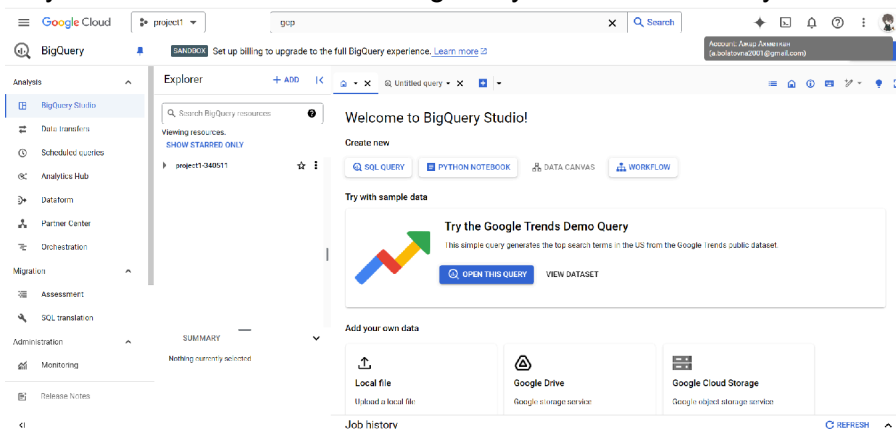
- How does Google Kubernetes Engine (GKE) simplify the management of containerized applications?



Google Kubernetes Engine simplifies the management of containerized applications by automating many tasks related to deployment and scaling. It handles load balancing, scaling, and monitoring, allowing

developers to focus more on building their apps instead of worrying about the infrastructure. Also, it automates updates and health checks, making the entire process more efficient.

- What advantages does Cloud Storage offer for data management? Cloud Storage offers several benefits for managing data. It is highly scalable - allowing users to store and access almost unlimited amounts of data without the fear of running out of space. It also provides high durability and availability, which automatically copying data across multiple locations. Another advantage is accessibility - users can reach their data from anywhere on the internet, making sharing and collaboration easy. Also it is cost-effective, with different pricing options that help businesses find what fits their budget and needs best.
- Why would a business choose BigQuery for their data analysis needs?



A business choose BigQuery for data analysis because it can quickly and efficiently manage large datasets. As a fully managed data warehouse, it enables real-time analytics and fast SQL queries, allowing companies to gain insights without handling complex infrastructure. Its integration with other Google Cloud services and support for machine learning make it an excellent option for businesses that want to use data for better decision making.

### Exercise 3: Creating and Managing Virtual Machines with Compute Engine

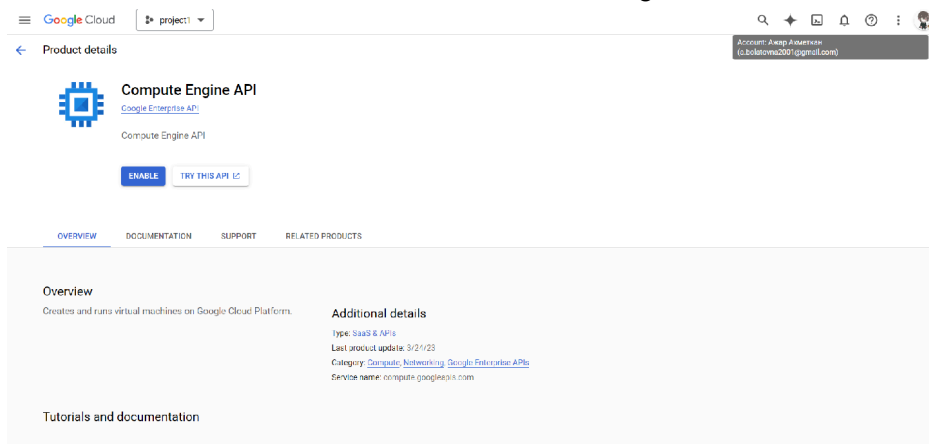
- **Objective:** Learn how to create, manage, and interact with virtual machines (VMs) using Compute Engine.
- **Steps:**
  - In the Google Cloud Console, navigate to Compute Engine and create a new VM instance.
  - Configure the VM with specific parameters, such as the machine type, region, and operating system.
  - Connect to the VM using SSH and install a basic web server (e.g., Apache or Nginx).
  - Stop, start, and delete the VM through the console.

- **Questions:**

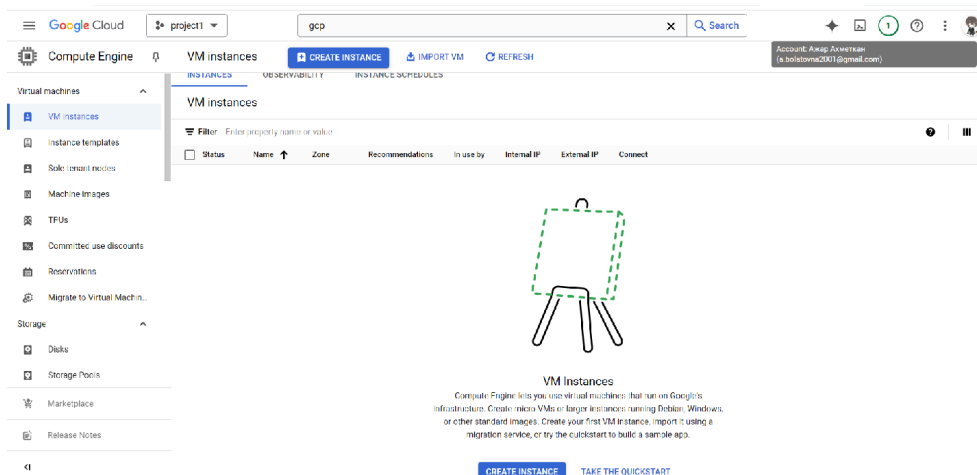
What steps did you follow to create the VM?

To create a VM instance in Google Cloud Console, I navigated to the Compute Engine

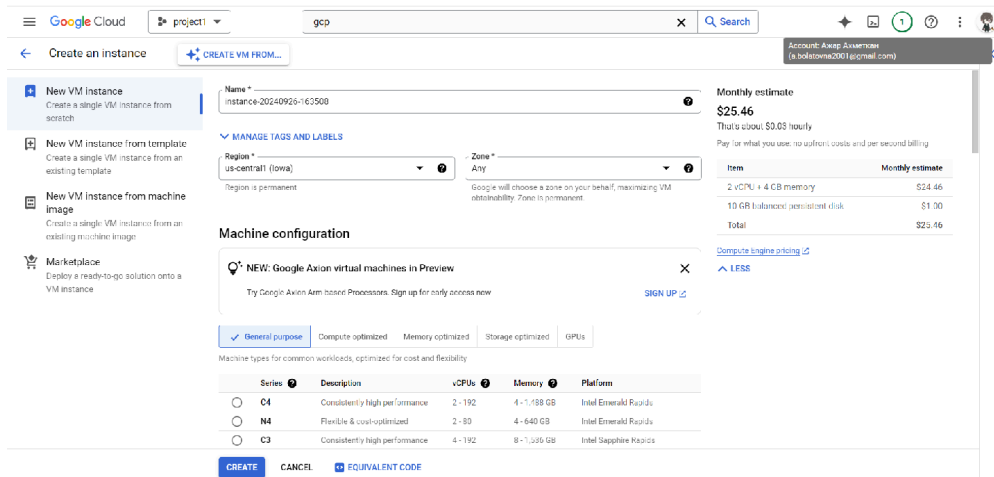
section and clicked button "enable" and add billing account.



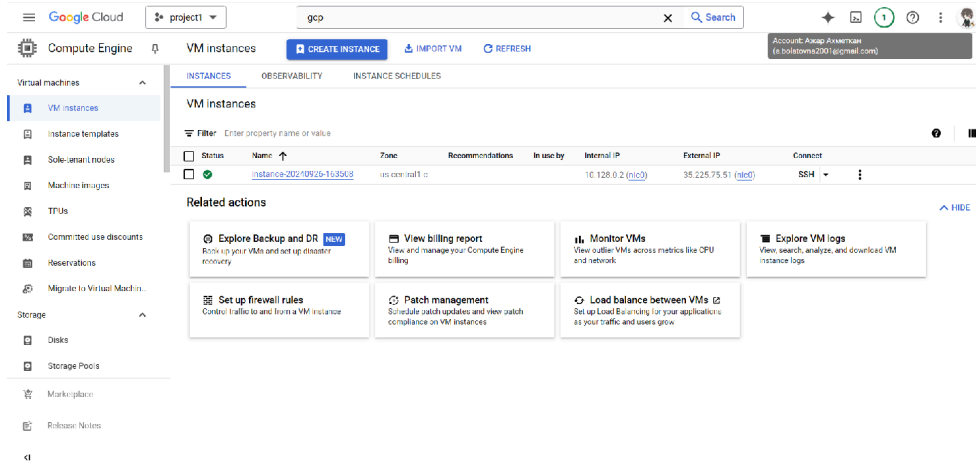
Next, clicked on "Create Instance."



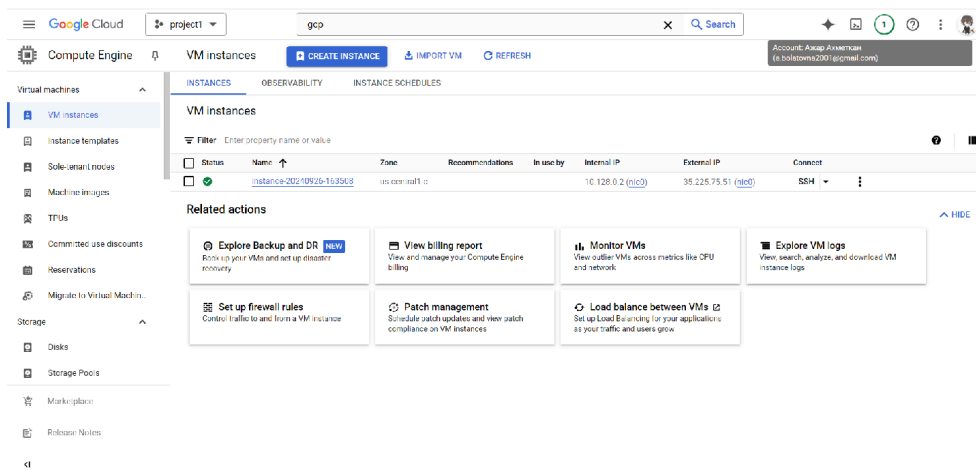
I then configured the VM by selecting the machine type, choosing the region, and picking the operating system I wanted to use. After confirming the settings, I clicked on the "Create" button to launch the VM.



Finally, below is my new VM instance



- How did you connect to the VM, and what commands did you use to install the web server?  
To connect to the VM, I used SSH. In the Google Cloud Console, I found the VM instance I created and clicked on the “SSH” button next to it.



This opened a terminal window where I could run commands. To install a basic web server Apache, I used

--sudo apt update

--sudo apt install apache2

```
ssh.cloud.google.com/v2/ssh/projects/project1-340511/zones/us-central1-c/instances/instance-20240926-1635...
ssh.cloud.google.com/v2/ssh/projects/project1-340511/zones/us-central1-c/instances/instance-20240926-...
SSH-in-browser
UPLOAD FILE
DOWNLOAD FILE

Linux instance-20240926-163508 6.1.0-25-cloud-amd64 #1 SMP PREEMPT_DYNAMIC Debian 6.1.106-3 (2024-08-26) x86_64

The programs included with the Debian GNU/Linux system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent
permitted by applicable law.

a_bolatovna2001@instance-20240926-163508:~$ sudo apt update
Get:1 file:/etc/apt/mirrors/debian.list Mirrorlist [30 B]
Get:5 file:/etc/apt/mirrors/debian-security.list Mirrorlist [39 B]
Get:7 https://packages.cloud.google.com/apt google-compute-engine-bookworm-stable InRelease [1321 B]
Get:2 https://deb.debian.org/debian bookworm InRelease [151 kB]
Get:8 https://packages.cloud.google.com/apt cloud-sdk-bookworm InRelease [1654 B]
Get:3 https://deb.debian.org/debian bookworm-updates InRelease [55.4 kB]
```

```
ssh.cloud.google.com/v2/ssh/projects/project1-340511/zones/us-central1-c/instances/instance-20240926-1635...
ssh.cloud.google.com/v2/ssh/projects/project1-340511/zones/us-central1-c/instances/instance-20240926-...
SSH-in-browser
UPLOAD FILE
DOWNLOAD FILE

Reading state information... Done
4 packages can be upgraded. Run 'apt list --upgradable' to see them.
a_bolatovna2001@instance-20240926-163508:~$ sudo apt install apache2
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  apache2-bin apache2-data apache2-utils libapr1 libaprutil1 libaprutil1-dbd-sqlite3
  libaprutil1-ldap libjansson4 liblua5.3-0 ssl-cert
Suggested packages:
  apache2-doc apache2-suexec-pristine | apache2-suexec-custom www-browser
The following NEW packages will be installed:
  apache2 apache2-bin apache2-data apache2-utils libapr1 libaprutil1 libaprutil1-dbd-sqlite3
  libaprutil1-ldap libjansson4 liblua5.3-0 ssl-cert
0 upgraded, 11 newly installed, 0 to remove and 4 not upgraded.
Need to get 2378 kB of archives.
After this operation, 8468 kB of additional disk space will be used.
Do you want to continue? [Y/n] Y
Get:1 file:/etc/apt/mirrors/debian.list Mirrorlist [30 B]
```

To install a basic web server Nginx, i used

--sudo apt update

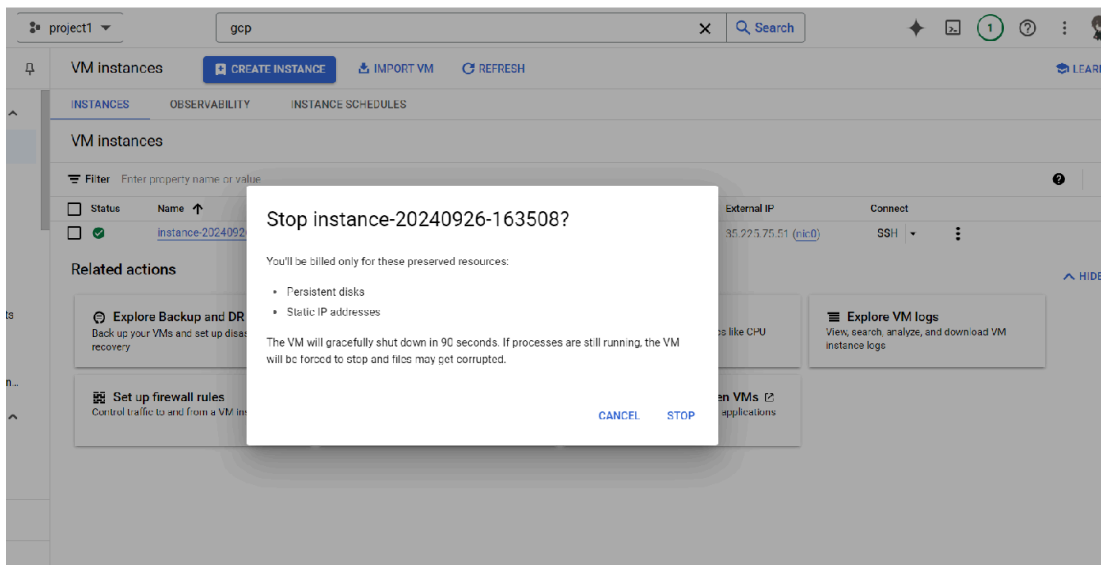
--sudo apt install nginx

```
systemd/system/apache-htcacheclean.service.
Processing triggers for man-db (2.11.2-2) ...
Processing triggers for libc-bin (2.36-9+deb12u8) ...
a_bolatovna2001@instance-20240926-163508:~$ sudo apt update
Get:1 file:/etc/apt/mirrors/debian.list Mirrorlist [30 B]
Get:2 file:/etc/apt/mirrors/debian-security.list Mirrorlist [39 B]
Hit:7 https://packages.cloud.google.com/apt google-compute-engine-bookworm-stable InRelease
Hit:8 https://packages.cloud.google.com/apt cloud-sdk-bookworm InRelease
Hit:3 https://deb.debian.org/debian bookworm InRelease
Hit:4 https://deb.debian.org/debian bookworm-updates InRelease
Hit:5 https://deb.debian.org/debian bookworm-backports InRelease
Hit:6 https://deb.debian.org/debian-security bookworm-security InRelease
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
4 packages can be upgraded. Run 'apt list --upgradable' to see them.
a_bolatovna2001@instance-20240926-163508:~$
```

```
Reading state information... Done
4 packages can be upgraded. Run 'apt list --upgradable' to see them.
a_bolatovna2001@instance-20240926-163508:~$ sudo apt install nginx
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  nginx-common
Suggested packages:
  fcgiwrap nginx-doc
The following NEW packages will be installed:
  nginx nginx-common
0 upgraded, 2 newly installed, 0 to remove and 4 not upgraded.
Need to get 640 kB of archives.
After this operation, 1696 kB of additional disk space will be used.
Do you want to continue? [Y/n] Y
```

- What happens to the VM and its data when it is stopped versus when it is deleted?

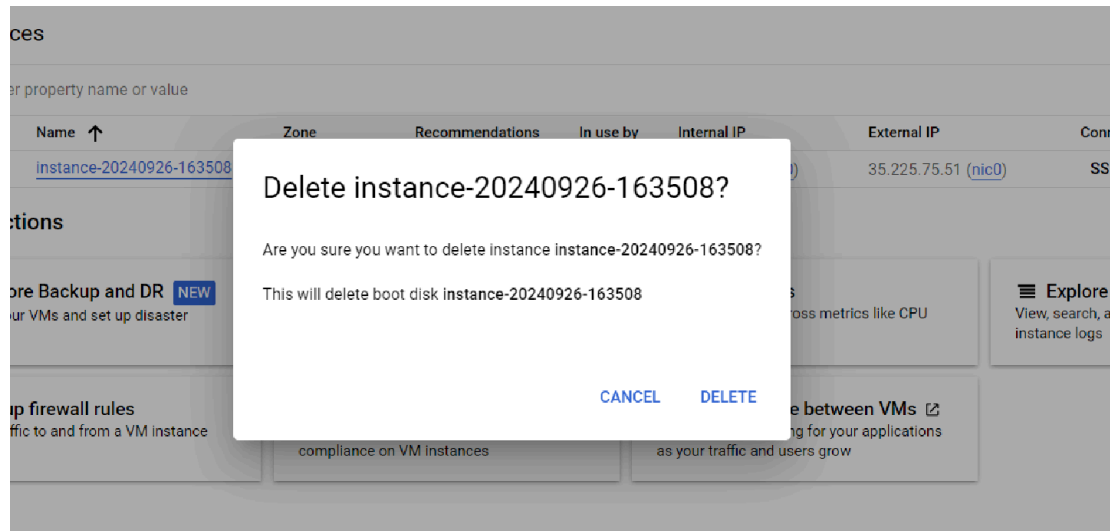
When the VM is stopped, it keeps all its data and settings, so it can be started again later. While the VM is stopped, You won't be charged for the resources while the VM is stopped, but you may still have charges for the persistent disk and any static IP addresses.



VM instance stopped



If the VM is deleted, all data on it is permanently lost, including any installed software and settings. However, if the data was saved on a separate persistent disk, that disk can be kept for future use.



Compute Engine lets you use virtual machines that run c  
infrastructure. Create micro-VMs or larger instances running [  
or other standard images. Create your first VM instance, im  
to build a sam

Instance deleted



CREATE INSTANCE

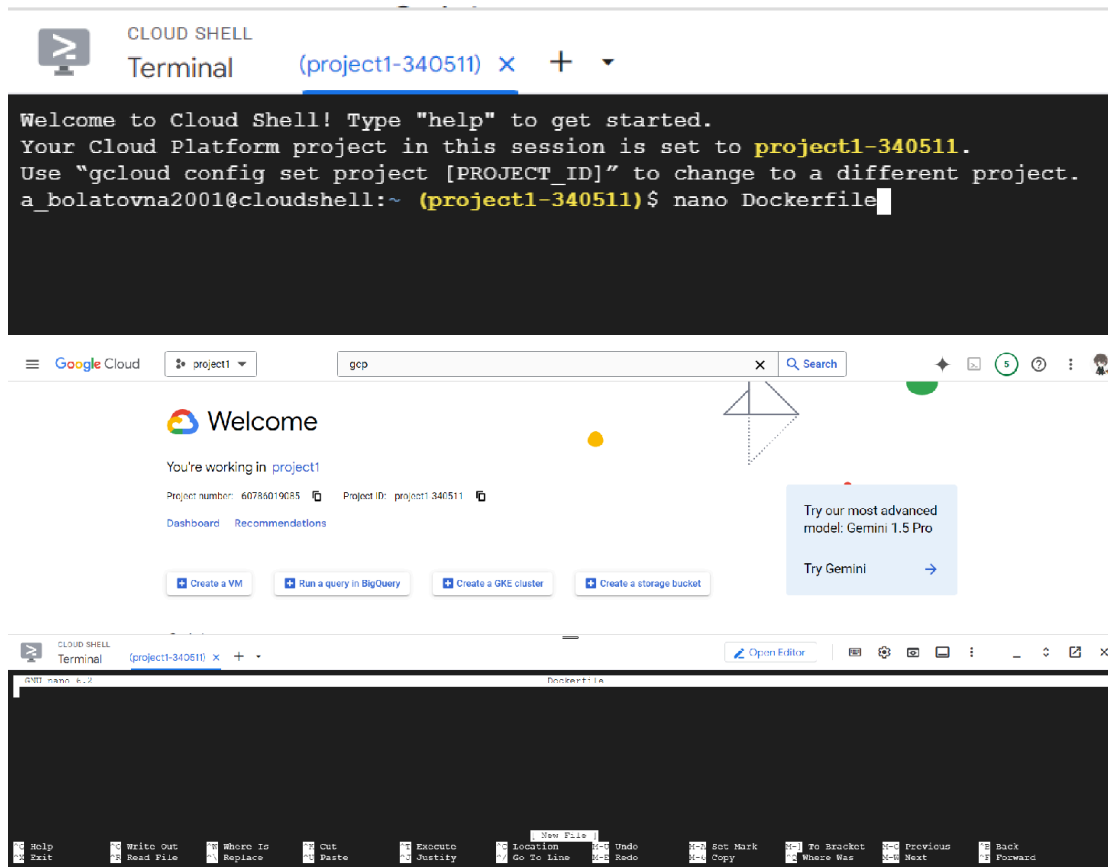
TAKE THE QUICKSTA

## Exercise 4: Deploying a Containerized Application on Google Kubernetes Engine (GKE)

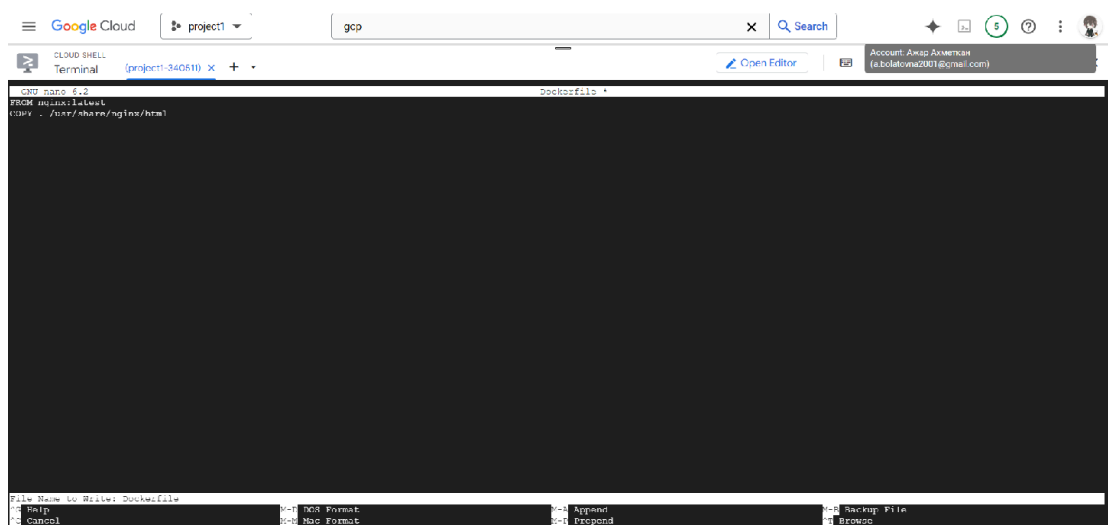
- **Objective:** Understand how to deploy and manage containerized applications using Google Kubernetes Engine.
- **Steps:**
  - Create a simple Docker container for a web application.
  - Push the container image to Google Container Registry (GCR).
  - Create a GKE cluster in Google Cloud Console.
  - Deploy the containerized application to the GKE cluster.
  - Expose the application to the internet and verify its accessibility.



- **Questions:**
  - How did you create and push the Docker container to GCR?  
I opened the Terminal, using Google Cloud Shell, open it from the Google Cloud Console and use a text editor nano in the terminal to create a Dockerfile.



Write the necessary instructions simple web application using Nginx in the Dockerfile and save it with pressing Ctrl + O




After saving, pressed Ctrl + X to exit the Nano editor and proceeded to build Docker image "azhar" by running the following command in the terminal:

```
--docker build -t azhar .
```

[illegible]

Next step was authenticate with Google Cloud, still in the terminal, I logged in to Google Cloud account and get window with password:

```
--gcloud auth login
```



## Sign in to the gcloud CLI

You are seeing this page because you ran the following command in the gcloud CLI from this or another machine. If this is not the case, close this tab.

```
gcloud auth login --no-launch-browser
```

Enter the following verification code in gcloud CLI on the machine you want to log into. This is a credential similar to your password and should not be shared with others.

```
4/0AQlEd8y7iYvP_AkStHwFL202i4-J0n10UAQ
nipScjayf4oJUzmyFkMWIz51HBwrQTMokGw
```

Copy

You can close this tab when you're done.

[illegible]

Using the terminal I tagged image for Google Container Registry and finally, pushed tagged image to Google Container Registry by commands below:

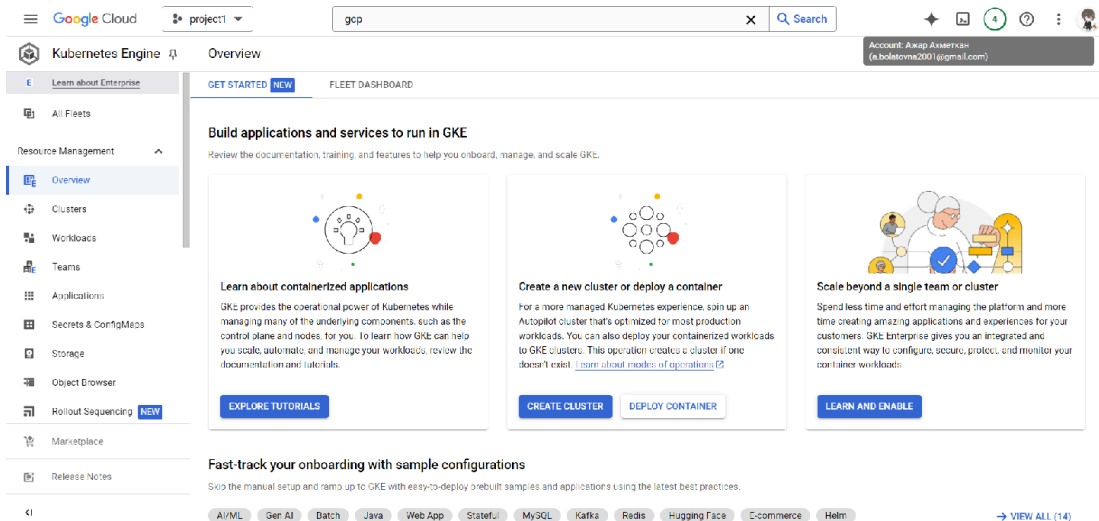
```
--docker tag azhar gcr.io/340511/azhar
```

```
--docker push gcr.io/340511/azhar
```

```
Once finished, enter the verification code provided in your browser: 4/0AQlEd8y7iYvP_AkStHWFI

You are now logged in as [a.bolatovna2001@gmail.com].
Your current project is [project1-340511]. You can change this setting by running:
$ gcloud config set project PROJECT_ID
a_bolatovna2001@cloudshell:~ (project1-340511)$ docker tag azhar gcr.io/340511/azhar
a_bolatovna2001@cloudshell:~ (project1-340511)$ docker push gcr.io/340511/azhar
Using default tag: latest
The push refers to repository [gcr.io/340511/azhar]
41be875c576e: Preparing
11de3d47036d: Preparing
16907864a2d0: Preparing
2bdf51597158: Preparing
0fc6bb94eec5: Preparing
eda13eb24d4c: Waiting
67796e30ff04: Waiting
8e2ab384f3bf: Waiting
```

- What steps were involved in setting up the GKE cluster?  
To set up the GKE cluster, I navigated to the Google Cloud Console and selected Kubernetes Engine. I clicked on “Create Cluster”



configured the settings, such as choosing the cluster type, specifying the number of nodes, and selecting the region.

Google Cloud project1 gcp Search

Account: Anap Arsenian (a.bolashina200@gmail.com)

### Create an Autopilot cluster

- Cluster basics**  
Set up basics for your cluster
- Fleet registration**  
Manage multiple clusters together
- Networking**  
Define applications communication in the cluster
- Advanced settings**  
Review additional options
- Review and create**  
Review all settings and create your cluster

#### Cluster basics

Create an Autopilot cluster by specifying a name and region. After the cluster is created, you can deploy your workload through Kubernetes and we'll take care of the rest, including:

- ✓ **Nodes:** Automated node provisioning, scaling, and maintenance
- ✓ **Networking:** VPC-native traffic routing for public or private clusters
- ✓ **Security:** Shielded GKE Nodes and Workload Identity
- ✓ **Telemetry:** Cloud Operations logging and monitoring

Name: autopilot-cluster-1

Region: us-central1

Cluster names must start with a lowercase letter, followed by up to 39 lowercase letters, numbers, or hyphens. They can't end with a hyphen. You cannot change the cluster's name once it's created.

The regional location in which your cluster's control plane and nodes are located. You cannot change the cluster's region once it's created.

[NEXT: FLEET REGISTRATION](#) [RESET SETTINGS](#)

[CREATE](#) [CANCEL](#) [Equivalent: REST](#) or [COMMAND LINE](#)

After reviewing the settings, I clicked “Create” to launch the cluster.

Google Cloud project1 gcp Search

Account: Anap Arsenian (a.bolashina200@gmail.com)

### Kubernetes Engine

[CREATE](#) [DEPLOY](#) [REFRESH](#)

[Learn about Enterprise](#)

All Fleets

Resource Management

- Overview
- Clusters**
- Workloads
- Teams
- Applications
- Secrets & ConfigMaps
- Storage
- Object Browser
- Rollout Sequencing **NEW**
- Marketplace
- Release Notes

Kubernetes clusters

[OVERVIEW](#) [OBSERVABILITY](#) [COST OPTIMIZATION](#)

Filter: Enter property name or value

Status	Name	Location	Mode	Number of nodes	Total vCPUs	Total memory	Notifications	Labels
✓	autopilot-cluster-1	us-central1	Autopilot	0	0	0 GB		

- How did you verify that your application was successfully deployed and accessible?

Open terminal Google Cloud Shell and I connected to GKE cluster

*gcloud container clusters get-credentials autopilot-cluster-1 --region us-central1 --project project1-340511*

After running the command, I checked the current context:

```
kubectx config current-context:
```

Used the command

```
kubectl create deployment azhardeployment --image=gcr.io/340511/azhar
```

to create a new deployment

To see the status of your deployments, I runned:

```
kubectl get deployments
```

Next, I checked the service associated with application to see how it's exposed to the internet.

I identified that it is ClusterIP: This type is only accessible within the cluster and won't work for external access. So i decided create new service with LoadBalancer.

## Check again services and get kubernetes2

So my external IP is 34.41.112.201 and can use to access application.

## Exercise 5: Storing and Accessing Data in Google Cloud Storage

- **Objective:** Learn how to store, manage, and access data using Google Cloud Storage.
- **Steps:**
  - Create a new Cloud Storage bucket in the Google Cloud Console.
  - Upload various types of files (e.g., text, images, videos) to the bucket.
- Set access permissions for the bucket and test public and private access to the files.
- Use the Cloud Console to download, move, and delete files in the bucket.
- **Questions:**
  - How do you create a Cloud Storage bucket, and what options are available during setup?

To create a Cloud Storage bucket, I go to the Google Cloud Console and navigate to the Cloud Storage section. Then, I click on "Create" near buckets.

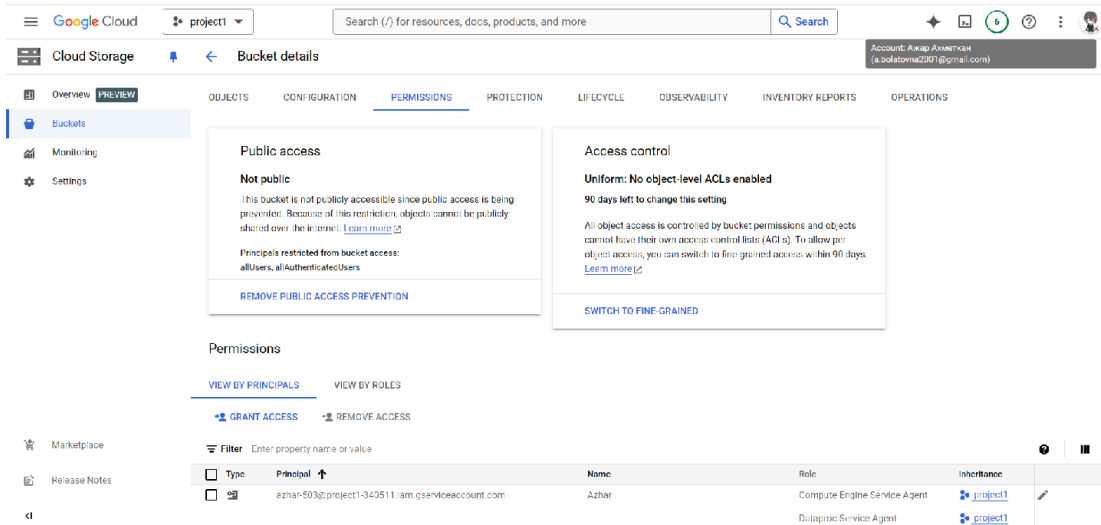
The screenshot shows the Google Cloud Console interface for the 'Buckets' section. The left sidebar contains navigation links for Overview, Buckets, Monitoring, and Settings. The main content area displays a table of existing buckets. The table has the following columns: Name, Created, Location type, Location, Default storage class, Last modified, and Public access. The buckets listed are:

Name	Created	Location type	Location	Default storage class	Last modified	Public access
azh-akb	Feb 7, 2022, 10:16:47 PM	Region	europe-west3	Standard	Feb 7, 2022, 10:16:47 PM	Not public
azhar-akhmed-khan-bucket	Feb 7, 2022, 7:42:14 PM	Multi-region	eu	Standard	Feb 7, 2022, 7:42:14 PM	Not public
dataproc-staging-europe-west3-56786d1...	Feb 7, 2022, 5:48:38 PM	Region	europe-west3	Standard	Feb 7, 2022, 5:48:38 PM	Subject to object ACL
dataproc-temp-europe-west3-50786c19...	Feb 7, 2022, 5:48:39 PM	Region	europe-west3	Standard	Feb 7, 2022, 5:48:39 PM	Subject to object ACL

Set up my bucket, choosing a unique name for my bucket, select a location and pick a storage class and decide on the access control method. Clicked create.

And got exercise\_5 bucket:

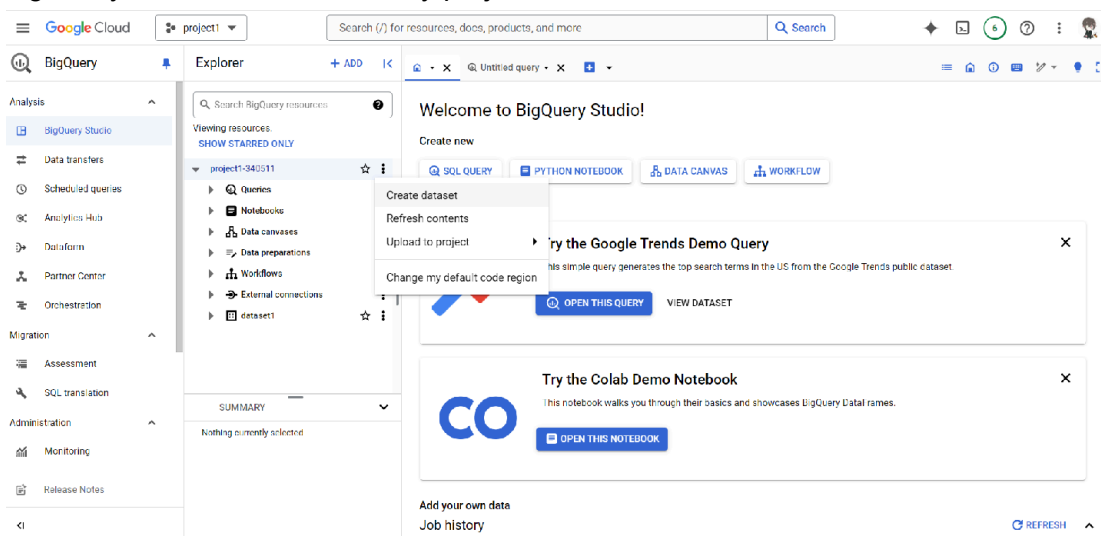
- What are the differences between setting a bucket to public versus private? Setting a bucket to public means anyone on the internet can access its files without needing authentication, which is great for sharing things widely. But a private bucket restricts access, so only people with the right permissions can view or download the files. This is important if I have sensitive data that shouldn't be accessible to everyone.
- How can you manage access permissions for individual files in a bucket? To manage access for individual files, I can select the file in the Cloud Console and go to the Permissions tab. From there, I can add or remove users, groups, or service accounts, and specify what roles they have, like Viewer or Editor. This way, I can control who can see or modify each file separately.



## Exercise 6: Analyzing Data with BigQuery

- **Objective:** Perform data analysis tasks using BigQuery.
- **Steps:**
  - Access BigQuery in the Google Cloud Console.
  - Create a dataset and table by importing a sample dataset provided by Google.
  - Write and execute SQL queries to perform basic data analysis, such as filtering, aggregation, and sorting.
  - Visualize the results using Google Data Studio or another visualization tool.
- **Questions:**
  - What steps did you take to create a dataset and table in BigQuery?

To create a dataset in BigQuery, I first logged into the Google Cloud Console and went to the BigQuery section. I clicked on my project, then hit "Create Dataset."



I picked a name "cloud\_dataset" for the dataset and set up the necessary options.



### Create dataset

Project ID \*

project1-340511

CHANGE

Dataset ID \*

cloud\_dataset

Letters, numbers, and underscores allowed

Location type ?

☐ Region

Specify a region to colocate your datasets with other Google Cloud services.

☒ Multi-region

Allow BigQuery to select a region within a group to achieve higher quota limits.

Multi-region \*

US (multiple regions in United States)

External Dataset

The selected region supports the following external dataset types: Cloud Spanner

☐ Link to an external dataset ?

Default table expiration

☐ Enable table expiration ?

Default maximum table age

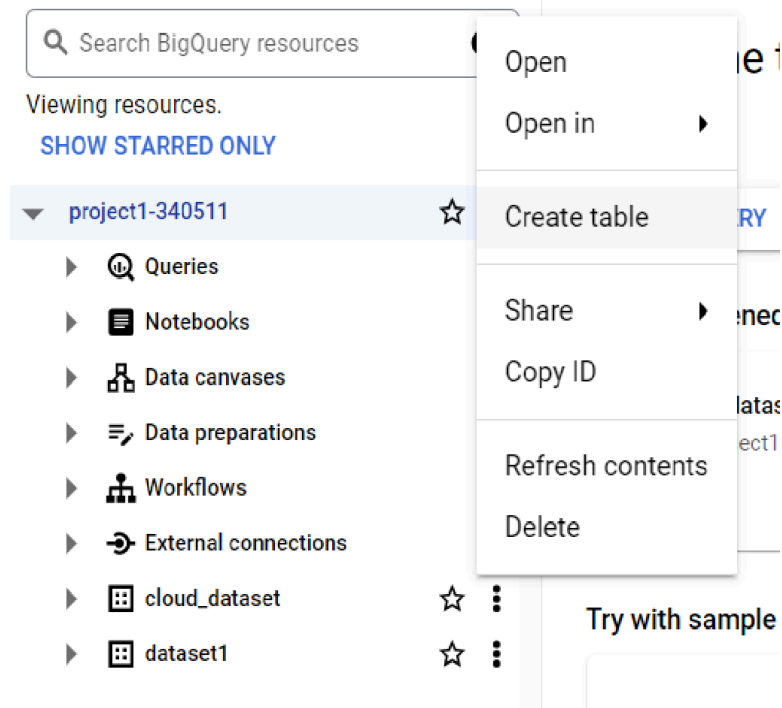
Days

Tags

CREATE DATASET

CANCEL

After that, I created a table "forexercise6" by clicking "Create Table,"



#### Schema

☒ Edit as text

1	Field name *	Type *	Mode		
	id	INTEGER	NULLABLE	Description	
2	Field name *	Type *	Mode	Max length	Description
	name	STRING	NULLABLE	100	
3	Field name *	Type *	Mode	Max length	Description
	surname	STRING	NULLABLE	200	
4	Field name *	Type *	Mode	Max length	Description
	course	STRING	NULLABLE	200	
5	+ ADD FIELD				

**CREATE TABLE** CANCEL

- How did you write and execute SQL queries in BigQuery?

I wrote SQL queries in the BigQuery editor. After selecting my dataset and table, I typed out the queries I wanted to run and clicked the "Run" button. I used commands to filter the data, do some aggregations, and sort the results based on what I needed to analyze.  
average performance score for each department:

Google Cloud project1 Search (/) for resources, docs, products, and more

BigQuery Explorer + ADD

Search BigQuery resources

Viewing resources. SHOW STARRED ONLY

project1-340511

Queries

Notebooks

Data canvases

Data preparations

Workflows

External connections

cloud\_dataset

exercise6

foreverexercise6

SHOW MORE

SUMMARY

Nothing currently selected

Untitled query

```
1 SELECT Department, AVG(Performance_Score) AS Average_Performance_Score
2 FROM `project1-340511.cloud_dataset.exercise6`
3 GROUP BY Department
4 ORDER BY Average_Performance_Score DESC;
```

Query results

JOB INFORMATION RESULTS CHART JSON EXECUTION DETAILS EXECUTION GRAPH

Row	Department	Average_Performance_Score
1	Engineering	3.02026869660...
2	Operations	3.036797748306...
3	IT	2.999999999999...
4	HR	2.997354014598...
5	Customer Support	2.997121266642...
6	Sales	2.99385851159...
7	Legal	2.981741320381...
8	Finance	2.981150714285...
9	Marketing	2.98062639087...

Results per page: 50 1 - 9 of 9

the employee satisfaction scores for employees who have not resigned:

Google Cloud project1 Search (/) for resources, docs, products, and more

BigQuery Explorer + ADD

Search BigQuery resources

Viewing resources. SHOW STARRED ONLY

project1-340511

Queries

Notebooks

Data canvases

Data preparations

Workflows

External connections

cloud\_dataset

exercise6

foreverexercise6

SHOW MORE

SUMMARY

Nothing currently selected

Untitled query

```
1 SELECT Employee_ID, Employee_Satisfaction_Score
2 FROM `project1-340511.cloud_dataset.exercise6`
3 WHERE Resigned = FALSE
4 ORDER BY Employee_Satisfaction_Score DESC;
```

Query results

JOB INFORMATION RESULTS CHART JSON EXECUTION DETAILS EXECUTION GRAPH

Row	Employee_ID	Employee_Satisfaction_Score
1	8869	5.0
2	5791	5.0
3	22552	5.0
4	35800	5.0
5	19222	5.0
6	95726	5.0
7	15708	5.0
8	3343	5.0
9	92885	5.0
10	79578	5.0
11	80318	5.0
12	75094	5.0
13	66298	5.0

Results per page: 50 1 - 50 of 59990

Employees with Overtime Hours:

Google Cloud project1 Search (/) for resources, docs, products, and more

BigQuery Explorer + ADD

Search BigQuery resources

Viewing resources. SHOW STARRED ONLY

project1-340511

Queries

Notebooks

Data canvases

Data preparations

Workflows

External connections

cloud\_dataset

exercise6

foreverexercise6

SHOW MORE

SUMMARY

Nothing currently selected

Untitled query

```
1 SELECT Employee_ID, Overtime_Hours
2 FROM `project1-340511.cloud_dataset.exercise6`
3 WHERE Overtime_Hours > 10
4 ORDER BY Overtime_Hours DESC;
```

Query results

JOB INFORMATION RESULTS CHART JSON EXECUTION DETAILS EXECUTION GRAPH

Row	Employee_ID	Overtime_Hours
1	3272	29
2	3601	29
3	4906	29
4	7102	29
5	7558	29
6	8745	29
7	8818	29
8	13414	29
9	13476	29
10	14073	29
11	14830	29
12	16673	29
13	18678	29

Results per page: 50 1 - 50 of 63415

- What insights were you able to derive from the data analysis?

From my analysis, I found a few key insights. First, the average performance score was highest in the IT department, suggesting they might have better support or engagement. Second, employees with more training hours had higher satisfaction scores, showing that training is important for keeping employees happy. Finally, many employees working significant overtime were in high-demand roles, which might indicate they are overloaded. This suggests that the company should consider balancing workloads to prevent burnout.